



MetaQA legacy gpt-oss, 200-per-hop: context and edge compression

Source: artifact_all/metaqa-vanilla-test; values match the gpt-oss column in the main result table.



HRG-Proposed-triple keeps the evidence context at 1.75% of BFS for this gpt-oss run.